

# COMP4318/5318 Assignment 2: Image Classification

Group number: A2 group 66 , SID1: 530311278 , SID2: 500312012

## Setup and dependencies

Please use this section to list and set up all your required libraries/dependencies and your plotting environment.

```
In [90]: import numpy as np
import pandas as pd
import random
import tensorflow as tf
import seaborn as sns
import keras_tuner as kt
# Import our deep learning libraries
from tensorflow import keras
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Make the notebook's output stable across runs
random.seed(42)
np.random.seed(42)
tf.random.set_seed(42)
keras.utils.set_random_seed(42)
keras.backend.clear_session()

%matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt

# Set the font size of axis titles and tick labels to make the chart more beautiful and easy to read.
mpl.rcParams['axes', labelsizes=14]
mpl.rcParams['xtick', labelsizes=12]
mpl.rcParams['ytick', labelsizes=12]
```

## 1. Data loading, exploration, and preprocessing

```
In [95]: # Load the dataset training and test sets as numpy arrays

X_train = np.load('Assignment2Data/X_train.npy')
y_train = np.load('Assignment2Data/y_train.npy')
X_test = np.load('Assignment2Data/X_test.npy')
y_test = np.load('Assignment2Data/y_test.npy')
```

### Examples of preprocessed data

Check out the shape of the data:

```
In [98]: print(f"Shape of X_train_full: {X_train.shape}")
print(f"Shape of y_train_full: {y_train.shape}")
print(f"Shape of X_test: {X_test.shape}")
print(f"Shape of y_test: {y_test.shape}")
```

```
Shape of X_train_full: (32000, 28, 28, 3)
Shape of y_train_full: (32000,)
Shape of X_test: (8000, 28, 28, 3)
Shape of y_test: (8000,)
```

Check the number of categories to see if they are balanced:

```
In [101]: # Training set
labels, counts = np.unique(y_train, return_counts=True)
for l, c in zip(labels, counts):
    print(f"Class {l}: {c} samples")
```

```
Class 0: 3490 samples
Class 1: 3431 samples
Class 2: 3505 samples
Class 3: 3656 samples
Class 4: 2950 samples
Class 5: 4290 samples
Class 6: 2728 samples
Class 7: 3253 samples
Class 8: 4697 samples
```

```
In [103... # Test set
labels, counts = np.unique(y_test, return_counts=True)
for l, c in zip(labels, counts):
    print(f"Class {l}: {c} samples")
```

```
Class 0: 873 samples
Class 1: 858 samples
Class 2: 877 samples
Class 3: 914 samples
Class 4: 737 samples
Class 5: 1072 samples
Class 6: 682 samples
Class 7: 813 samples
Class 8: 1174 samples
```

## Analysis of Pixel Value Distribution

### 1. Pixel values are mostly concentrated between 120 and 230

- This indicates that the images are generally bright, with few dark regions (e.g., 0–50, typically representing shadows or background).
- There are very few extreme values at 0 or 255, meaning the dataset **does not contain “completely black” or “completely white” images**. This suggests good overall image quality.

### 2. The distribution is right-skewed but smooth

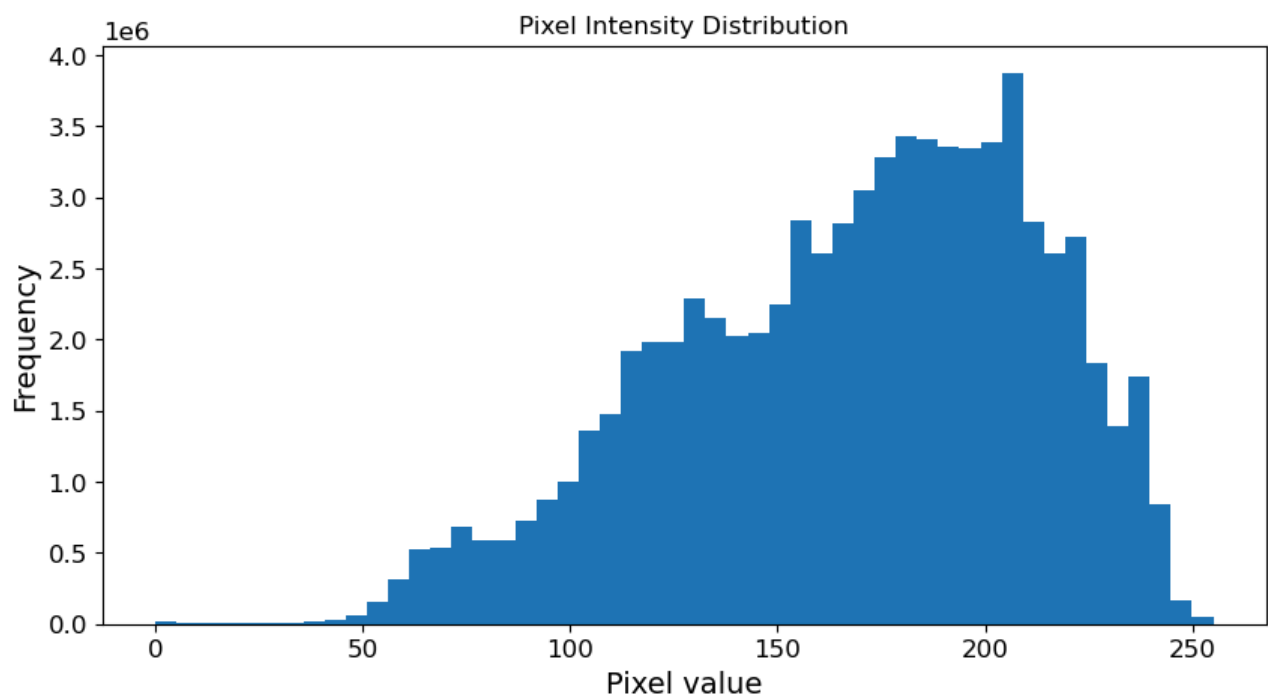
- Most images appear bright backgrounds, which is a common characteristic of histopathology tissue slides.
- The histogram shows a generally increasing trend followed by a gradual decline, **without abrupt spikes or bimodal patterns**. This suggests consistent background coloration across the dataset.

### 3. Pixel values are continuous but not yet normalized

- The raw pixel values range from **0 to 255**. Before training a neural network model, **normalization (i.e., dividing by 255) is necessary** to ensure numerical stability during learning.

```
In [107... # Flatten all image pixels into a one-dimensional array
# and plot a histogram of the pixel intensity distribution
```

```
plt.figure(figsize=(10, 5))
plt.hist(X_train.ravel(), bins=50)
plt.title("Pixel Intensity Distribution")
plt.xlabel("Pixel value")
plt.ylabel("Frequency")
plt.show()
```



**Grayscale conversion:** take the average across the RGB channels to convert color images to grayscale.

- The original `X_train` shape is `(num_samples, 28, 28, 3)`, and after conversion it becomes `(num_samples, 28, 28, 1)`.
- We initially attempted to train the model using grayscale images to simplify the input structure and speed up training. However, **experimental results showed that performance was inferior compared to using color images**. Therefore, grayscale images are now only used for visual comparison with the original color images.

```
In [110... X_train_gray = np.mean(X_train, axis=-1, keepdims=True) # shape: (num_samples, 28, 28, 1)
X_test_gray = np.mean(X_test, axis=-1, keepdims=True)
```

## Sample image

- Visualize 3 sample images per label in both original color and grayscale format. This helps assess whether the images are **clear, centered, and exhibit distinguishable features**. The comparison also provides insight into whether grayscale retains sufficient structural information.

```
In [113... # Get all unique class labels
labels = np.unique(y_train)
# Number of images to display per class
samples_per_class = 3


# Create a subplot grid: one row per class, two columns per sample (color + grayscale)
fig, axes = plt.subplots(
    len(labels), samples_per_class * 2,
    figsize=(samples_per_class * 4, len(labels) * 2.2)
)


for row, label in enumerate(labels):
    # Find all indices in y_train that match the current label
    label_indices = np.where(y_train == label)[0]
    # Select the first N samples
    selected_indices = label_indices[:samples_per_class]

    for i, idx in enumerate(selected_indices):
        # Left: display color image
        ax_color = axes[row, i * 2]
        ax_color.imshow(X_train[idx]) # Original color image
        if i == 0:
            ax_color.set_ylabel(f"Label: {label}", fontsize=11)
            ax_color.set_title(f"Label: {y_train[idx]}", fontsize=9)
            ax_color.axis('off')

        # Right: display corresponding grayscale image
        ax_gray = axes[row, i * 2 + 1]
        ax_gray.imshow(X_train_gray[idx].squeeze(), cmap='gray') # Grayscale image
        ax_gray.set_title(f"Label: {y_train[idx]}", fontsize=9)
        ax_gray.axis('off')

plt.tight_layout()
plt.show()
```





## Data preprocessing

```
In [116... # MLP + CNN
# Scale the data to the range 0-1
X_train_scale = X_train / 255.0
X_test_scale = X_test / 255.0

In [118... # Random Forest cannot handle multidimensional image inputs directly,
# so we reshape each (28, 28, 3) image into a 1D feature vector.

X_train_rf = X_train_scale.reshape((X_train_scale.shape[0], -1))
X_test_rf = X_test_scale.reshape((X_test_scale.shape[0], -1))
```

## Creating a validation set

To improve the training efficiency of neural networks, we introduced a third dataset — **the validation set**. Although this may reduce evaluation stability, it is more efficient than using k-fold cross-validation, which involves significantly higher computational cost. **Using a single split for validation is more practical for monitoring training and tuning hyperparameters.**

- During training, we use the validation set to evaluate model performance:
  - **To prevent overfitting** (e.g., by using early stopping to halt training early),
  - **To select the best hyperparameters.**
- **Prevent information leakage.**
  - The test set must remain completely independent from the entire training process and should only be used for final evaluation. If the test set participates in training, the model may learn from it, leading to overestimated performance and losing its purpose of assessing generalization.

```
In [121... # MLP + CNN: Create validation set from the training set
X_train_nn, X_valid, y_train_nn, y_valid = train_test_split(
    X_train_scale, y_train, train_size=0.9, random_state=42, stratify=y_train
)
```

## 2. Algorithm design and setup

### Algorithm of choice from first six weeks of course - Random Forest

#### Baseline - Random Forest

- The main function of the baseline model is to provide a performance reference point for subsequent model optimization (such as hyperparameter tuning).
- We try to keep the model simple (**no parameters are adjusted/parameters are kept as simple as possible**)

```
In [70]: # Baseline - RF
# Initialize the model with default hyperparameters (set random_state for reproducibility)
rf = RandomForestClassifier(random_state=42)

# Train
rf.fit(X_train_rf, y_train)

# Predict on the test set
y_pred = rf.predict(X_test_rf)
acc_base_rf = accuracy_score(y_test, y_pred)

# Print evaluation metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Classification Report:					
	precision	recall	f1-score	support	
0	0.82	0.98	0.89	873	
1	0.99	0.90	0.94	858	
2	0.58	0.32	0.42	877	
3	0.73	0.74	0.74	914	
4	0.54	0.65	0.59	737	
5	0.62	0.73	0.67	1072	
6	0.49	0.12	0.19	682	
7	0.49	0.46	0.47	813	
8	0.52	0.74	0.61	1174	
accuracy			0.65	8000	
macro avg	0.64	0.63	0.61	8000	
weighted avg	0.64	0.65	0.63	8000	

## Fully connected neural network

### Baseline - MLP

- Built a simple network using default settings: no learning rate tuning, no Dropout, no EarlyStopping
- 2 Dense layers: 128 → 64
- Fixed architecture and hyperparameters
- Optimizer: Adam (default learning rate)

```
In [126... def build_mlp_baseline():
    model = keras.models.Sequential()

    # Input layer + Flatten: convert (28, 28, 3) image to 1D vector
    model.add(keras.layers.Input(shape=(28, 28, 3)))
    model.add(keras.layers.Flatten())

    # First Dense layer -> 128 units + relu
    model.add(keras.layers.Dense(128, activation='relu'))

    # Second Dense layer -> 64 units + relu
    model.add(keras.layers.Dense(64, activation='relu'))

    # Output layer -> 9 units (number of classes) + softmax (probability output)
    model.add(keras.layers.Dense(9, activation='softmax'))

    model.compile(
        optimizer = "Adam",
        loss = 'sparse_categorical_crossentropy',
        metrics = ['accuracy']
    )

    return model
```

```
In [128... # Train the baseline

baseline_model_mlp = build_mlp_baseline()

baseline_model_mlp.fit(
    X_train_nn, y_train_nn,
    epochs = 10,
    validation_data = (X_valid, y_valid)
)
```

Epoch 1/10  
**900/900** ————— 2s 2ms/step - accuracy: 0.2344 - loss: 1.9929 - val\_accuracy: 0.3613 - val\_loss: 1.6765  
Epoch 2/10  
**900/900** ————— 2s 2ms/step - accuracy: 0.3565 - loss: 1.6838 - val\_accuracy: 0.4069 - val\_loss: 1.5973  
Epoch 3/10  
**900/900** ————— 1s 2ms/step - accuracy: 0.3956 - loss: 1.5921 - val\_accuracy: 0.4531 - val\_loss: 1.4810  
Epoch 4/10  
**900/900** ————— 2s 2ms/step - accuracy: 0.4431 - loss: 1.4865 - val\_accuracy: 0.4809 - val\_loss: 1.3821  
Epoch 5/10  
**900/900** ————— 1s 2ms/step - accuracy: 0.4715 - loss: 1.4215 - val\_accuracy: 0.4984 - val\_loss: 1.3595  
Epoch 6/10  
**900/900** ————— 2s 2ms/step - accuracy: 0.4857 - loss: 1.3852 - val\_accuracy: 0.5147 - val\_loss: 1.3257  
Epoch 7/10  
**900/900** ————— 1s 2ms/step - accuracy: 0.4928 - loss: 1.3663 - val\_accuracy: 0.5231 - val\_loss: 1.3097  
Epoch 8/10  
**900/900** ————— 2s 2ms/step - accuracy: 0.4995 - loss: 1.3455 - val\_accuracy: 0.5309 - val\_loss: 1.2821  
Epoch 9/10  
**900/900** ————— 1s 2ms/step - accuracy: 0.5032 - loss: 1.3314 - val\_accuracy: 0.5344 - val\_loss: 1.2675  
Epoch 10/10  
**900/900** ————— 2s 2ms/step - accuracy: 0.5081 - loss: 1.3149 - val\_accuracy: 0.5297 - val\_loss: 1.2570

Out[128... <keras.src.callbacks.history.History at 0x30e5a8830>

```
In [129... loss_base_mlp, acc_base_mlp = baseline_model_mlp.evaluate(X_test_scale, y_test)
print(f"Baseline Test Accuracy: {acc_base_mlp:.4f}")
```

250/250 ————— 0s 504us/step – accuracy: 0.5203 – loss: 1.2506  
Baseline Test Accuracy: 0.5279

```
In [130... y_pred = baseline_model_mlp.predict(X_test_scale)
y_pred_labels = np.argmax(y_pred, axis=1)

print(classification_report(y_test, y_pred_labels))
```

250/250 ————— 0s 368us/step

	precision	recall	f1-score	support
0	0.74	0.96	0.84	873
1	0.92	0.71	0.80	858
2	0.38	0.31	0.34	877
3	0.52	0.66	0.58	914
4	0.54	0.55	0.54	737
5	0.48	0.56	0.52	1072
6	0.33	0.03	0.06	682
7	0.23	0.14	0.17	813
8	0.43	0.65	0.52	1174
accuracy			0.53	8000
macro avg	0.51	0.51	0.49	8000
weighted avg	0.51	0.53	0.50	8000

- After training for 10 epochs, the model achieved an accuracy of about 50% on the test set.

## Algorithm Design - MLP

- **Input Layer:** Accepts 3-channel color images of shape (28, 28, 3) -> flattened to 1D.
- **Two hidden layers:**
  - Layer 1: 256–1024 (step=64)
  - Layer 2: 128–512 (step=64)
  - ReLU activation
- **Dropout layers** are applied after **each Dense layer (dropout rate between 0.2 and 0.5)** -> prevent overfitting.
  - **Excessive dropout** may harm model capacity, leading to underfitting.
  - **Large fluctuations** in validation performance may occur if the network becomes too unstable due to frequent dropout.
- **Output Layer:** 9-unit Dense layer + softmax activation multi-class classification
- **Hyperparameter tuning:**
  - units, dropout\_rate, and learning\_rate are all tunable.
  - learning rate [1e-4, 1e-2] + log-scale sampling.

```
In [133... def build_mlp(hp):
    model = keras.models.Sequential()
    # Input layer – accepts images with shape (28, 28, 3)
    model.add(keras.layers.Input(shape=(28, 28, 3)))

    # Flatten the 3D image into a 1D vector (28*28*3 = 2352)
    model.add(keras.layers.Flatten())

    # Define dropout rate as a tunable hyperparameter
    dropout_rate = hp.Float('dropout_rate', 0.2, 0.5, step=0.1)

    # First hidden layer – extract low-level features
    model.add(
        keras.layers.Dense(
            units=hp.Int('units_1', min_value=256, max_value=1024, step=64),
            activation='relu'
        )
    )
    # Dropout to prevent overfitting
    model.add(keras.layers.Dropout(dropout_rate))

    # Second hidden layer – fuse features into higher-level representation
    model.add(
        keras.layers.Dense(
            units=hp.Int('units_2', min_value=128, max_value=512, step=64),
            activation='relu'
        )
    )
    # Dropout to prevent overfitting
```

```

model.add(keras.layers.Dropout(dropout_rate))

# Output layer -> 9 units (number of classes) + softmax (probability output)
model.add(keras.layers.Dense(9, activation='softmax'))

# Instantiate optimiser and compile the model.
model.compile(
    optimizer=keras.optimizers.Adam(
        learning_rate=hp.Float('lr', 1e-4, 1e-2, sampling='log')
    ),
    loss='sparse_categorical_crossentropy', # suitable for integer class labels (not one-hot!)
    metrics=['accuracy']
)

return model

```

## Convolutional neural network

### Baseline - CNN

- Built a simple CNN model: no learning rate tuning, no Dropout, no EarlyStopping
- 2 convolutional blocks followed by one dense layer
- fixed architecture and hyperparameters

```

In [136... def build_cnn_baseline():
    model = keras.models.Sequential()

    # Input layer - accepts images with shape (28, 28, 3)
    model.add(keras.layers.Input(shape = (28, 28, 3)))

    # First block
    model.add(keras.layers.Conv2D(filters = 32, kernel_size = (3, 3), activation = 'relu'))
    model.add(keras.layers.MaxPooling2D(pool_size = (2, 2)))

    # Second block
    model.add(keras.layers.Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu'))
    model.add(keras.layers.MaxPooling2D(pool_size = (2, 2)))

    # Flatten
    model.add(keras.layers.Flatten())

    # Fully connected layer
    model.add(keras.layers.Dense(64, activation = 'relu'))

    # Output layer -> 9 units (number of classes) + softmax (probability output)
    model.add(keras.layers.Dense(9, activation = 'softmax'))

    model.compile(
        optimizer = "Adam",
        loss = 'sparse_categorical_crossentropy',
        metrics = ['accuracy']
    )

    return model

```

```

In [137... baseline_model_cnn = build_cnn_baseline()

baseline_model_cnn.fit(
    X_train_nn, y_train_nn,
    epochs = 10,
    validation_data = (X_valid, y_valid)
)

```



```

Epoch 1/10
900/900 ————— 5s 6ms/step - accuracy: 0.3580 - loss: 1.6963 - val_accuracy: 0.5959 - val_loss: 1.0724
Epoch 2/10
900/900 ————— 5s 6ms/step - accuracy: 0.6150 - loss: 1.0454 - val_accuracy: 0.6684 - val_loss: 0.9090
Epoch 3/10
900/900 ————— 5s 5ms/step - accuracy: 0.6680 - loss: 0.9121 - val_accuracy: 0.6953 - val_loss: 0.8379
Epoch 4/10
900/900 ————— 5s 6ms/step - accuracy: 0.7061 - loss: 0.8190 - val_accuracy: 0.7444 - val_loss: 0.7355
Epoch 5/10
900/900 ————— 5s 6ms/step - accuracy: 0.7274 - loss: 0.7643 - val_accuracy: 0.7541 - val_loss: 0.6962
Epoch 6/10
900/900 ————— 5s 5ms/step - accuracy: 0.7362 - loss: 0.7329 - val_accuracy: 0.7544 - val_loss: 0.6979
Epoch 7/10
900/900 ————— 5s 5ms/step - accuracy: 0.7460 - loss: 0.7018 - val_accuracy: 0.7584 - val_loss: 0.6869
Epoch 8/10
900/900 ————— 5s 6ms/step - accuracy: 0.7605 - loss: 0.6670 - val_accuracy: 0.7628 - val_loss: 0.6732
Epoch 9/10
900/900 ————— 5s 6ms/step - accuracy: 0.7691 - loss: 0.6373 - val_accuracy: 0.7700 - val_loss: 0.6555
Epoch 10/10
900/900 ————— 6s 6ms/step - accuracy: 0.7755 - loss: 0.6200 - val_accuracy: 0.7788 - val_loss: 0.6232

```

Out[137... <keras.src.callbacks.history.History at 0x314f68620>

```

In [138... loss_base_cnn, acc_base_cnn = baseline_model_cnn.evaluate(X_test_scale, y_test)
print(f"Baseline Test Accuracy: {acc_base_cnn:.4f}")

```

```

250/250 ————— 1s 3ms/step - accuracy: 0.7769 - loss: 0.6142
Baseline Test Accuracy: 0.7739

```

```

In [139... y_pred = baseline_model_cnn.predict(X_test_scale)
y_pred_labels = np.argmax(y_pred, axis=1)

print(classification_report(y_test, y_pred_labels))

```

```

250/250 ————— 1s 2ms/step

```

	precision	recall	f1-score	support
0	0.86	0.97	0.91	873
1	0.98	0.92	0.95	858
2	0.63	0.65	0.64	877
3	0.97	0.89	0.93	914
4	0.72	0.77	0.74	737
5	0.67	0.83	0.75	1072
6	0.66	0.62	0.64	682
7	0.59	0.50	0.54	813
8	0.85	0.75	0.80	1174
accuracy			0.77	8000
macro avg	0.77	0.77	0.77	8000
weighted avg	0.78	0.77	0.77	8000

## Algorithm Design - CNN

- The model includes **three convolutional layers**, with increasing filters, followed by a fully connected layer and output layer.
  - Input: (28, 28, 3)
  - Conv2D → MaxPooling2D → Dropout
  - Conv2D (filters × 2) → MaxPooling2D → Dropout
  - Conv2D (filters × 2)
  - Flatten
  - Dense (filters × 2) → Dropout
  - Dense(9) with softmax activation
- Key hyperparameters tuned:**
  - filters: Controls the number of feature detectors [32, 64, 128]
  - dropout\_rate: Controls the regularization strength [0.2 - 0.5]
  - learning\_rate: Tuned on a log scale [1e-4 - 1e-2]
- The core architecture of the baseline is retained, but the flexibility is enhanced → making it easier to optimize performance in Keras Tuner

```

In [169... def build_cnn(hp):
    model = keras.models.Sequential()
    # Input layer - accepts images with shape (28, 28, 3)
    model.add(keras.layers.Input(shape=(28, 28, 3)))

    # Hyperparameters to tune - filters and dropout rate
    filters = hp.Choice('filters', values=[32, 64, 128]) # Controls model capacity
    dropout_rate = hp.Float('dropout_rate', 0.2, 0.5, step=0.1) # Prevent overfitting

    # First block
    model.add(keras.layers.Conv2D(filters, kernel_size=(3, 3), activation='relu'))
    model.add(keras.layers.MaxPooling2D(pool_size=(2, 2)))

```

```

model.add(keras.layers.Dropout(dropout_rate))           # Prevent overfitting

# Second block
model.add(keras.layers.Conv2D(filters * 2, kernel_size=(3, 3), activation='relu'))
model.add(keras.layers.MaxPooling2D(pool_size=(2, 2)))
model.add(keras.layers.Dropout(dropout_rate))           # Prevent overfitting

# Third block - no pooling
model.add(keras.layers.Conv2D(filters * 2, kernel_size=(3, 3), activation='relu'))

# Flatten
model.add(keras.layers.Flatten())

# Fully connected layer
model.add(keras.layers.Dense(filters * 2, activation='relu'))
# Dropout after dense layer - prevents reliance on specific neurons
model.add(keras.layers.Dropout(dropout_rate))

# Output layer -> 9 units (number of classes) + softmax (probability output)
model.add(keras.layers.Dense(9, activation='softmax'))

model.compile(
    optimizer=keras.optimizers.Adam(
        learning_rate=hp.Float('lr', 1e-4, 1e-2, sampling='log')
    ),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

return model

```

### 3. Hyperparameter tuning

#### Algorithm of choice from first six weeks of course - Random Forest

- **GridSearchCV** already performs cross-validation internally (cv=5):
- It will **automatically split** the provided X\_train\_rf, y\_train into 5 folds;
- Each fold is internally divided into **train/validation sets** to evaluate the performance of different parameter combinations.

```

In [104... # Define the parameter grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'max_features': ['sqrt', 'log2']
}

# Grid search
grid_search_rf = GridSearchCV(rf, param_grid, cv=5)
grid_search_rf.fit(X_train_rf, y_train)

# Print the best parameters
best_params_rf = grid_search_rf.best_params_
print("Best hyperparameters:", best_params_rf)

```

Best hyperparameters: {'max\_depth': 30, 'max\_features': 'sqrt', 'min\_samples\_split': 5, 'n\_estimators': 200}

#### Visualization

```

In [106... # Convert the grid search results into a DataFrame
cv_results = pd.DataFrame(grid_search_rf.cv_results_)

cv_summary = cv_results[[
    'param_n_estimators',
    'param_max_depth',
    'param_min_samples_split',
    'param_max_features',
    'mean_test_score',
    'mean_fit_time'
]]
cv_sorted = cv_summary.sort_values(by='mean_test_score', ascending=False)

# Print all parameter combinations - training time + accuracy
print(cv_sorted)

# Export - plotting / analysis in the report
cv_sorted.to_csv('rf_gridsearch_results.csv', index=False)

```

	param_n_estimators	param_max_depth	param_min_samples_split	\
41	200	30	5	
44	200	30	10	
38	200	30	2	
23	200	20	5	
26	200	20	10	
20	200	20	2	
40	100	30	5	
47	200	30	2	
43	100	30	10	
53	200	30	10	
32	200	20	5	
50	200	30	5	
35	200	20	10	
29	200	20	2	
37	100	30	2	
22	100	20	5	
19	100	20	2	
25	100	20	10	
52	100	30	10	
34	100	20	10	
46	100	30	2	
31	100	20	5	
49	100	30	5	
28	100	20	2	
42	50	30	10	
24	50	20	10	
39	50	30	5	
21	50	20	5	
18	50	20	2	
36	50	30	2	
51	50	30	10	
33	50	20	10	
48	50	30	5	
30	50	20	5	
45	50	30	2	
27	50	20	2	
5	200	10	5	
2	200	10	2	
8	200	10	10	
4	100	10	5	
7	100	10	10	
1	100	10	2	
3	50	10	5	
0	50	10	2	
6	50	10	10	
14	200	10	5	
17	200	10	10	
11	200	10	2	
13	100	10	5	
10	100	10	2	
16	100	10	10	
12	50	10	5	
15	50	10	10	
9	50	10	2	

	param_max_features	mean_test_score	mean_fit_time
41	sqrt	0.649875	100.701498
44	sqrt	0.648438	98.852448
38	sqrt	0.648344	102.206669
23	sqrt	0.647969	97.191994
26	sqrt	0.646563	94.642976
20	sqrt	0.645844	98.557626
40	sqrt	0.638281	50.370415
47	log2	0.638281	26.814444
43	sqrt	0.638000	49.346964
53	log2	0.637656	24.391276
32	log2	0.637406	24.600942
50	log2	0.637312	25.593390
35	log2	0.636656	23.699626
29	log2	0.636500	25.183181
37	sqrt	0.635938	50.404275
22	sqrt	0.635125	48.426561
19	sqrt	0.634844	49.522002
25	sqrt	0.634375	47.036766
52	log2	0.627500	12.496703
34	log2	0.626125	11.847130
46	log2	0.626063	13.893486
31	log2	0.624500	12.348277
49	log2	0.623563	13.800292
28	log2	0.622937	12.593124
42	sqrt	0.620250	25.940749
24	sqrt	0.618125	23.637237
39	sqrt	0.616781	25.900678
21	sqrt	0.614469	24.408040

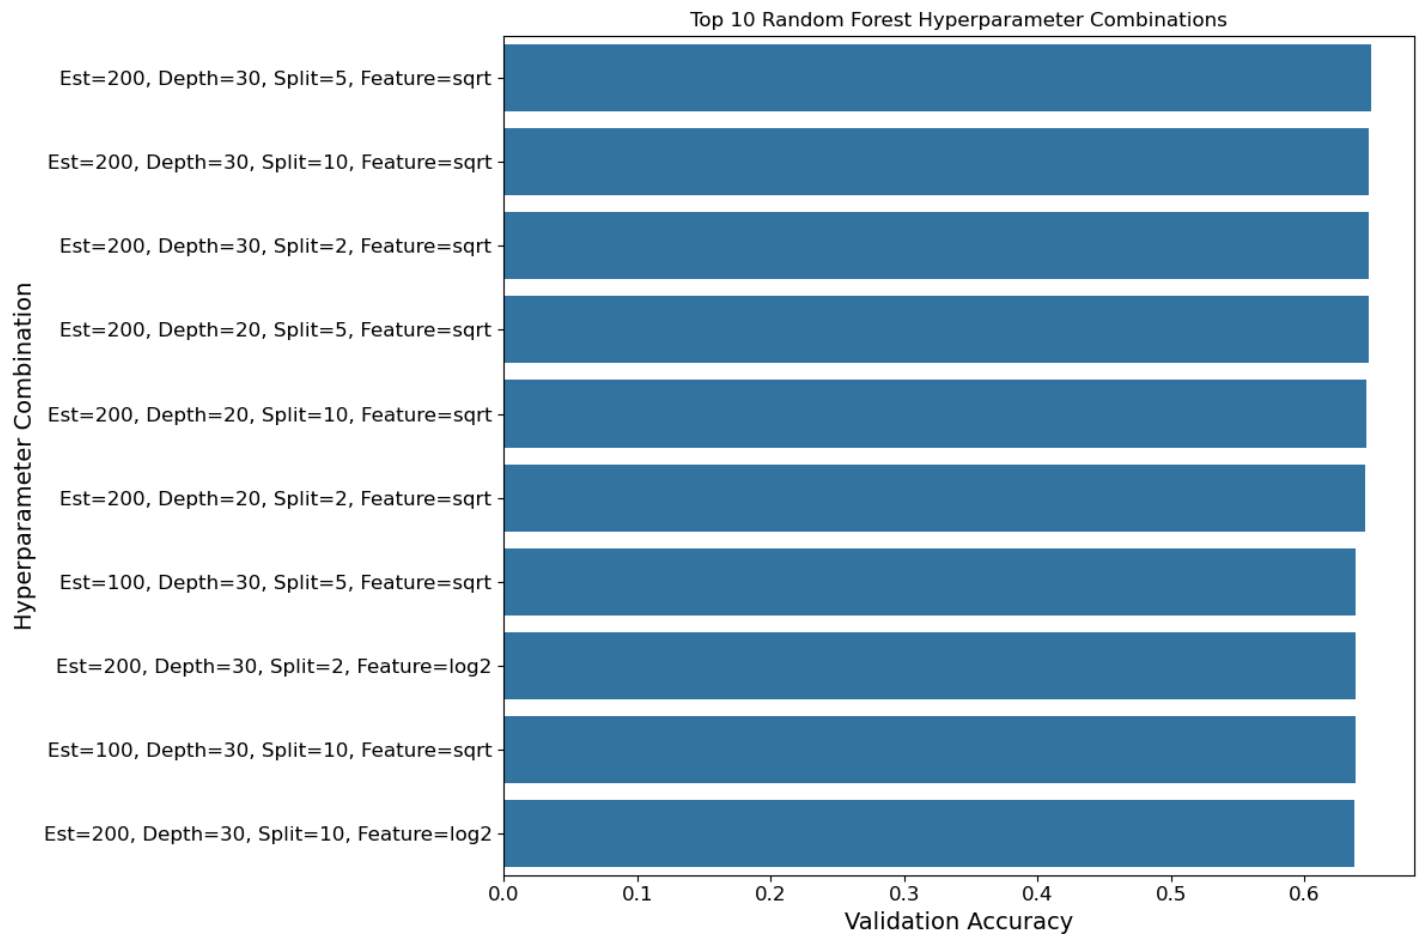
18	sqrt	0.613125	24.726600
36	sqrt	0.610188	25.274157
51	log2	0.608750	6.098683
33	log2	0.608563	5.918097
48	log2	0.605312	7.089012
30	log2	0.604344	6.219662
45	log2	0.601344	7.064115
27	log2	0.599969	6.304334
5	sqrt	0.598781	60.620036
2	sqrt	0.597844	60.488921
8	sqrt	0.597469	60.157498
4	sqrt	0.593094	30.299843
7	sqrt	0.592844	30.079155
1	sqrt	0.590281	31.041808
3	sqrt	0.582937	14.953385
0	sqrt	0.582469	16.248459
6	sqrt	0.581187	15.224582
14	log2	0.581094	14.999165
17	log2	0.580844	14.800021
11	log2	0.579719	15.064560
13	log2	0.575937	7.542744
10	log2	0.574625	7.452049
16	log2	0.574281	7.424692
12	log2	0.568188	3.824771
15	log2	0.566156	3.758625
9	log2	0.564875	3.946650

```
In [106... rf_df = pd.read_csv('rf_gridsearch_results.csv')

# Create tag
rf_df['param_combo'] = (
    'Est=' + rf_df['param_n_estimators'].astype(str) +
    ', Depth=' + rf_df['param_max_depth'].astype(str) +
    ', Split=' + rf_df['param_min_samples_split'].astype(str) +
    ', Feature=' + rf_df['param_max_features'].astype(str)
)
```

```
In [106... # Sort by accuracy - DESC
rf_sorted = rf_df.sort_values(by='mean_test_score', ascending=False)

# Plot
plt.figure(figsize=(12, 8))
sns.barplot(x='mean_test_score', y='param_combo', data=rf_sorted.head(10))
plt.xlabel('Validation Accuracy')
plt.ylabel('Hyperparameter Combination')
plt.title('Top 10 Random Forest Hyperparameter Combinations')
plt.tight_layout()
plt.show()
```



## Fully connected neural network

```
In [824... # Use Random Search
tuner_mlp = kt.RandomSearch(
    hypermodel = build_mlp,
    objective = 'val_accuracy',
    max_trials = 30,                # Search 30 different hyperparameter combinations
    executions_per_trial = 1,       # Run each combination once
    overwrite = True,              # Overwrite previous results
    directory = 'keras_tuning_results', # Results folder
    project_name = 'mlp'
)

# Stop if no improvement for 5 consecutive epochs - reduce the computing time
early_stopping = EarlyStopping(
    monitor = 'val_accuracy',
    patience = 5,
    restore_best_weights = True
)

# Print a summary of the search
tuner_mlp.search_space_summary()

# Hyperparameter search
tuner_mlp.search(
    X_train_nn, y_train_nn,
    epochs = 30,
    validation_data = (X_valid, y_valid),
    callbacks=[early_stopping]
)
```

Trial 30 Complete [00h 00m 39s]  
val\_accuracy: 0.1471875011920929

Best val\_accuracy So Far: 0.5512499809265137  
Total elapsed time: 00h 22m 46s

## Visualization

```
In [825... # Retrieve the top 10 best performing trials from the MLP tuner
mlp_trials = tuner_mlp.oracle.get_best_trials(num_trials=10)

# Format the results
mlp_records = []
for trial in mlp_trials:
```

```

record = trial.hyperparameters.values.copy()
record['score'] = trial.score # this is the final val_accuracy
best_epoch = trial.metrics.get_best_step('val_accuracy') # epoch with best performance
record['best_epoch'] = best_epoch + 1 # +1 to make epoch indexing start from 1
mlp_records.append(record)

# Convert the records into a DataFrame
mlp_df = pd.DataFrame(mlp_records)
print(mlp_df)

# Save the results
mlp_df.to_csv('mlp_tuner_results.csv', index=False)

```

	dropout_rate	units_1	units_2	lr	score	best_epoch
0	0.2	448	320	0.000104	0.551250	28
1	0.2	704	256	0.000338	0.523438	9
2	0.2	832	256	0.000267	0.522500	14
3	0.3	576	320	0.000122	0.517812	22
4	0.3	896	128	0.000277	0.511250	17
5	0.3	448	320	0.000171	0.504375	12
6	0.2	256	128	0.000335	0.491250	10
7	0.3	256	128	0.000275	0.474687	10
8	0.2	640	448	0.000454	0.465000	3
9	0.4	256	128	0.000122	0.438750	12

```

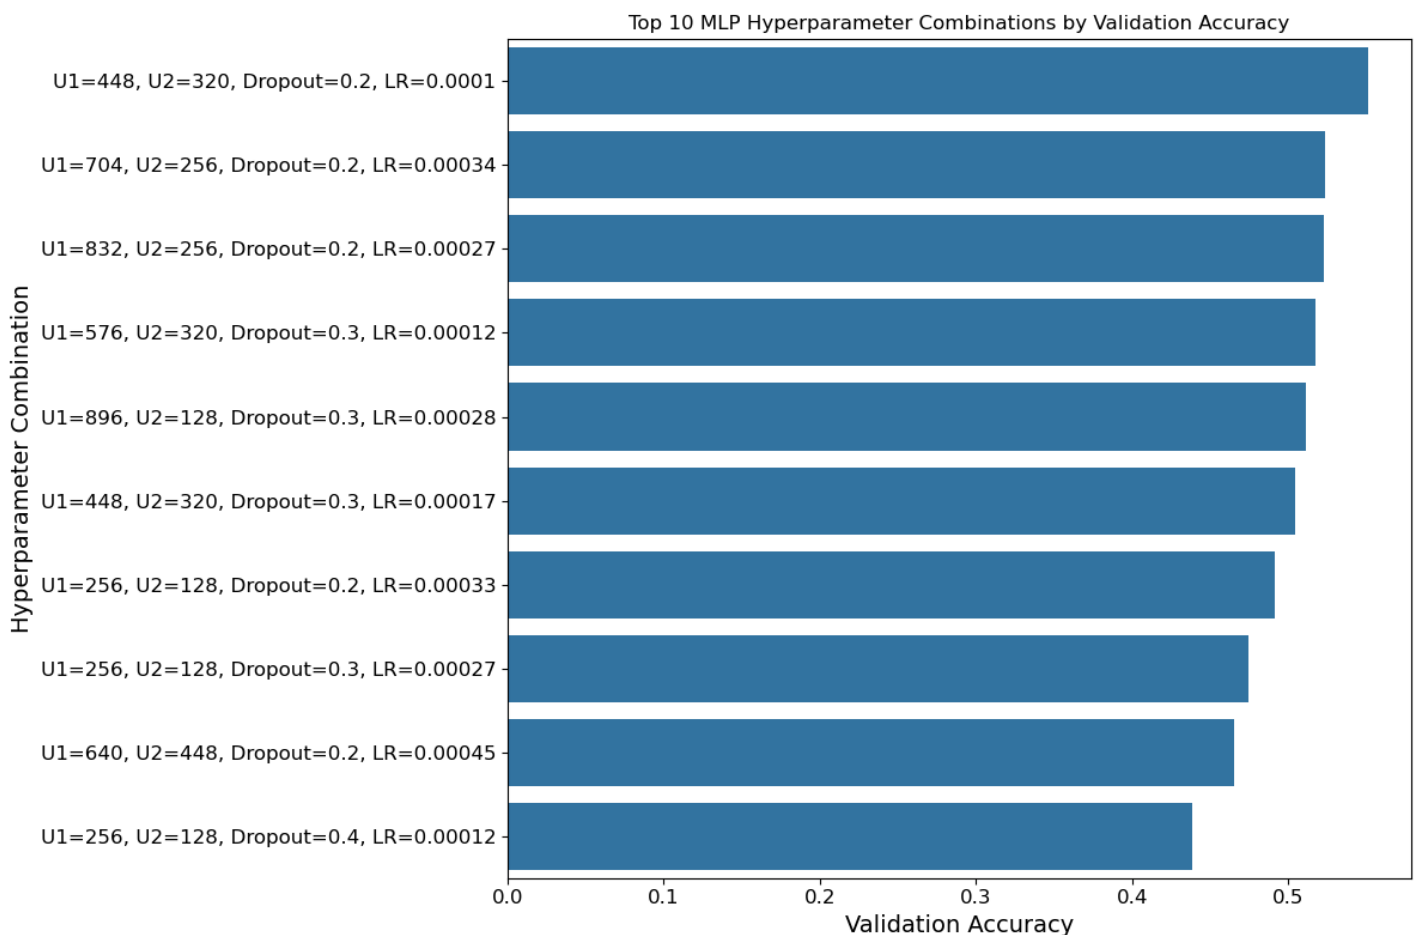
In [834... mlp_df = pd.read_csv('mlp_tuner_results.csv')

# Create tag
mlp_df['param_combo'] = (
    'U1=' + mlp_df['units_1'].astype(str) +
    ', U2=' + mlp_df['units_2'].astype(str) +
    ', Dropout=' + mlp_df['dropout_rate'].astype(str) +
    ', LR=' + mlp_df['lr'].round(5).astype(str)
)


# Sort by accuracy - DESC
mlp_df_sorted = mlp_df.sort_values('score', ascending=False)

# barplot (top 10 hyperparameter combinations)
plt.figure(figsize=(12, 8))
sns.barplot(x='score', y='param_combo', data=mlp_df_sorted.head(10))
plt.xlabel('Validation Accuracy')
plt.ylabel('Hyperparameter Combination')
plt.title('Top 10 MLP Hyperparameter Combinations by Validation Accuracy')
plt.tight_layout()
plt.show()


```




```
In [836... # Dropout Rate vs Validation Accuracy (bar)
plt.figure(figsize=(8, 5))
sns.barplot(x='dropout_rate', y='score', data=mlp_df)
plt.title('MLP: Dropout Rate vs Validation Accuracy')
plt.xlabel('Dropout Rate')
plt.ylabel('Validation Accuracy')
plt.tight_layout()
plt.show()
```




```
In [838... # Dense Layer Units vs Accuracy (scatter)
for units in ['units_1', 'units_2']:
    plt.figure(figsize=(8, 5))
    sns.scatterplot(x=units, y='score', data=mlp_df)
    plt.title(f'MLP: {units} vs Validation Accuracy')
    plt.xlabel(units)
    plt.ylabel('Validation Accuracy')
    plt.tight_layout()
    plt.show()
```





```
In [840... # Learning Rate vs Accuracy (log scale)
plt.figure(figsize=(8, 5))
sns.scatterplot(x='lr', y='score', data=mlp_df)
plt.xscale('log')
plt.title('MLP: Learning Rate vs Validation Accuracy')
plt.xlabel('Learning Rate (log scale)')
plt.ylabel('Validation Accuracy')
plt.tight_layout()
plt.show()
```



### Validation Model with Best Hyperparameters

To evaluate the training dynamics of the MLP model under the best-performing hyperparameter combination, we manually rebuild the model using:

- units\_1 = 448
- units\_2 = 320
- dropout\_rate = 0.2
- learning\_rate = 0.000104

We train the model for a fixed number of 28 epochs, and **plot the accuracy and loss curves** on both the training and validation sets. This helps us:



- Visualize the learning process
- Detect overfitting by comparing the gap between training and validation performance
- Decide whether to adjust the number of training epochs

This step ensures the chosen architecture and parameters generalize well

```
In [103... def valid_mlp():
    model = keras.models.Sequential()
    model.add(keras.layers.Input(shape = (28, 28, 3)))
    model.add(keras.layers.Flatten())

    # Fixed parameters
    model.add(keras.layers.Dense(448, activation = 'relu'))
    model.add(keras.layers.Dropout(0.2))

    model.add(keras.layers.Dense(320, activation = 'relu'))
    model.add(keras.layers.Dropout(0.2))

    model.add(keras.layers.Dense(9, activation = 'softmax'))

    model.compile(
        optimizer = keras.optimizers.Adam(learning_rate = 0.000104),
        loss = 'sparse_categorical_crossentropy',
        metrics = ['accuracy']
    )



    return model
```

```
In [103... # Fixed epoch = 28
valid = valid_mlp()
history_mlp = valid.fit(
    X_train_nn, y_train_nn,
    epochs = 28,
    validation_data = (X_valid, y_valid)
)

# Accuracy
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history_mlp.history['accuracy'], label='Train Accuracy')
plt.plot(history_mlp.history['val_accuracy'], label='Val Accuracy')
plt.title('MLP - Training & Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

# Loss
plt.subplot(1, 2, 2)
plt.plot(history_mlp.history['loss'], label='Train Loss')
plt.plot(history_mlp.history['val_loss'], label='Val Loss')
plt.title('MLP - Training & Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()
```

Epoch 1/28  
**900/900** ————— 5s 5ms/step - accuracy: 0.2102 - loss: 2.0610 - val\_accuracy: 0.3319 - val\_loss: 1.7335  
Epoch 2/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.3328 - loss: 1.7473 - val\_accuracy: 0.4244 - val\_loss: 1.6226  
Epoch 3/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.3655 - loss: 1.6722 - val\_accuracy: 0.4363 - val\_loss: 1.5599  
Epoch 4/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.3922 - loss: 1.6142 - val\_accuracy: 0.4525 - val\_loss: 1.5138  
Epoch 5/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4119 - loss: 1.5735 - val\_accuracy: 0.4672 - val\_loss: 1.4783  
Epoch 6/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4300 - loss: 1.5354 - val\_accuracy: 0.4809 - val\_loss: 1.4324  
Epoch 7/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4442 - loss: 1.5056 - val\_accuracy: 0.4794 - val\_loss: 1.4145  
Epoch 8/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4529 - loss: 1.4777 - val\_accuracy: 0.4831 - val\_loss: 1.3857  
Epoch 9/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4617 - loss: 1.4631 - val\_accuracy: 0.4888 - val\_loss: 1.3839  
Epoch 10/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4605 - loss: 1.4474 - val\_accuracy: 0.4941 - val\_loss: 1.3726  
Epoch 11/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4662 - loss: 1.4349 - val\_accuracy: 0.4884 - val\_loss: 1.3691  
Epoch 12/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4701 - loss: 1.4228 - val\_accuracy: 0.4916 - val\_loss: 1.3525  
Epoch 13/28  
**900/900** ————— 5s 5ms/step - accuracy: 0.4791 - loss: 1.4091 - val\_accuracy: 0.4847 - val\_loss: 1.3617  
Epoch 14/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4770 - loss: 1.4035 - val\_accuracy: 0.5028 - val\_loss: 1.3362  
Epoch 15/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4816 - loss: 1.3957 - val\_accuracy: 0.5003 - val\_loss: 1.3280  
Epoch 16/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4855 - loss: 1.3877 - val\_accuracy: 0.5072 - val\_loss: 1.3100  
Epoch 17/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4914 - loss: 1.3748 - val\_accuracy: 0.5147 - val\_loss: 1.3110  
Epoch 18/28  
**900/900** ————— 4s 5ms/step - accuracy: 0.4905 - loss: 1.3719 - val\_accuracy: 0.4994 - val\_loss: 1.3405  
Epoch 19/28  
**900/900** ————— 4s 5ms/step - accuracy: 0.4948 - loss: 1.3606 - val\_accuracy: 0.5125 - val\_loss: 1.2860  
Epoch 20/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.4968 - loss: 1.3550 - val\_accuracy: 0.5134 - val\_loss: 1.2849  
Epoch 21/28  
**900/900** ————— 4s 5ms/step - accuracy: 0.5023 - loss: 1.3446 - val\_accuracy: 0.5181 - val\_loss: 1.2828  
Epoch 22/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.5007 - loss: 1.3436 - val\_accuracy: 0.5213 - val\_loss: 1.2745  
Epoch 23/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.5032 - loss: 1.3289 - val\_accuracy: 0.5319 - val\_loss: 1.2754  
Epoch 24/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.5054 - loss: 1.3291 - val\_accuracy: 0.5253 - val\_loss: 1.2633  
Epoch 25/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.5082 - loss: 1.3197 - val\_accuracy: 0.5378 - val\_loss: 1.2467  
Epoch 26/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.5049 - loss: 1.3186 - val\_accuracy: 0.5319 - val\_loss: 1.2515  
Epoch 27/28  
**900/900** ————— 4s 5ms/step - accuracy: 0.5128 - loss: 1.3020 - val\_accuracy: 0.5337 - val\_loss: 1.2382  
Epoch 28/28  
**900/900** ————— 4s 4ms/step - accuracy: 0.5162 - loss: 1.3009 - val\_accuracy: 0.5406 - val\_loss: 1.2297



## Convolutional neural network

```
In [919... # Use Random Search
tuner_cnn = kt.RandomSearch(
    hypermodel = build_cnn,
    objective = 'val_accuracy',
    max_trials = 10,
    executions_per_trial = 1,
    # Search 10 different hyperparameter combinations
    # Run each combination once
```

```
    overwrite = True,                # Overwrite previous results
    directory = 'keras_tuning_results', # Results folder
    project_name = 'cnn'
)

# Early stopping has already defined in the previous(mlp) part
# Print a summary of the search
tuner_cnn.search_space_summary()

# Hyperparameter search
tuner_cnn.search(
    X_train_nn, y_train_nn,
    epochs = 30,
    validation_data = (X_valid, y_valid),
    callbacks = [early_stopping]
)

tuner_cnn.results_summary()
```

```
Trial 10 Complete [00h 01m 34s]
val_accuracy: 0.7799999713897705

Best val_accuracy So Far: 0.8853124976158142
Total elapsed time: 01h 13m 59s
Results summary
Results in keras_tuning_results/cnn
Showing 10 best trials
Objective(name="val_accuracy", direction="max")
```

```
Trial 04 summary
Hyperparameters:
filters: 128
dropout_rate: 0.30000000000000004
lr: 0.00013612460334701279
Score: 0.8853124976158142
```

```
Trial 05 summary
Hyperparameters:
filters: 64
dropout_rate: 0.30000000000000004
lr: 0.0002458321717641159
Score: 0.8678125143051147
```

```
Trial 02 summary
Hyperparameters:
filters: 128
dropout_rate: 0.30000000000000004
lr: 0.00011641658137192351
Score: 0.8668749928474426
```

```
Trial 01 summary
Hyperparameters:
filters: 128
dropout_rate: 0.30000000000000004
lr: 0.0006529927317508998
Score: 0.8615624904632568
```

```
Trial 07 summary
Hyperparameters:
filters: 64
dropout_rate: 0.2
lr: 0.00018974491572285046
Score: 0.8578125238418579
```

```
Trial 06 summary
Hyperparameters:
filters: 128
dropout_rate: 0.4
lr: 0.0007506269481542165
Score: 0.8518750071525574
```

```
Trial 09 summary
Hyperparameters:
filters: 32
dropout_rate: 0.30000000000000004
lr: 0.0013960407115272237
Score: 0.7799999713897705
```

```
Trial 00 summary
Hyperparameters:
filters: 32
dropout_rate: 0.4
lr: 0.0012158303534721207
Score: 0.7584375143051147
```

```
Trial 03 summary
Hyperparameters:
filters: 64
dropout_rate: 0.2
lr: 0.006172427728586157
Score: 0.4568749964237213
```

```
Trial 08 summary
Hyperparameters:
filters: 64
dropout_rate: 0.2
lr: 0.008779360094579483
Score: 0.14687499403953552
```

## Visualization

```
In [921]: # Retrieve the top 10 best performing trials from the CNN tuner
cnn_trials = tuner_cnn.oracle.get_best_trials(num_trials=10)
```

```

# Format the results
cnn_records = []
for trial in cnn_trials:
    record = trial.hyperparameters.values.copy()
    record['score'] = trial.score
    record['best_epoch'] = trial.metrics.get_best_step('val_accuracy') + 1 # this is the final val accuracy
                                # +1 to make epoch indexing start from 1
    cnn_records.append(record)

# Convert the records into a DataFrame
cnn_df = pd.DataFrame(cnn_records)
print(cnn_df)

# Save the results
cnn_df.to_csv('cnn_tuner_results.csv', index=False)

```

	filters	dropout_rate	lr	score	best_epoch
0	128	0.3	0.000136	0.885312	30
1	64	0.3	0.000246	0.867813	29
2	128	0.3	0.000116	0.866875	29
3	128	0.3	0.000653	0.861562	25
4	64	0.2	0.000190	0.857813	30
5	128	0.4	0.000751	0.851875	30
6	32	0.3	0.001396	0.780000	13
7	32	0.4	0.001216	0.758438	25
8	64	0.2	0.006172	0.456875	12
9	64	0.2	0.008779	0.146875	1

In [932... `cnn_df = pd.read_csv('cnn_tuner_results.csv')`

```

# Create tag
cnn_df['param_combo'] = (
    'Filter=' + cnn_df['filters'].astype(str) +
    ', Dropout=' + cnn_df['dropout_rate'].astype(str) +
    ', LR=' + cnn_df['lr'].round(5).astype(str)
)

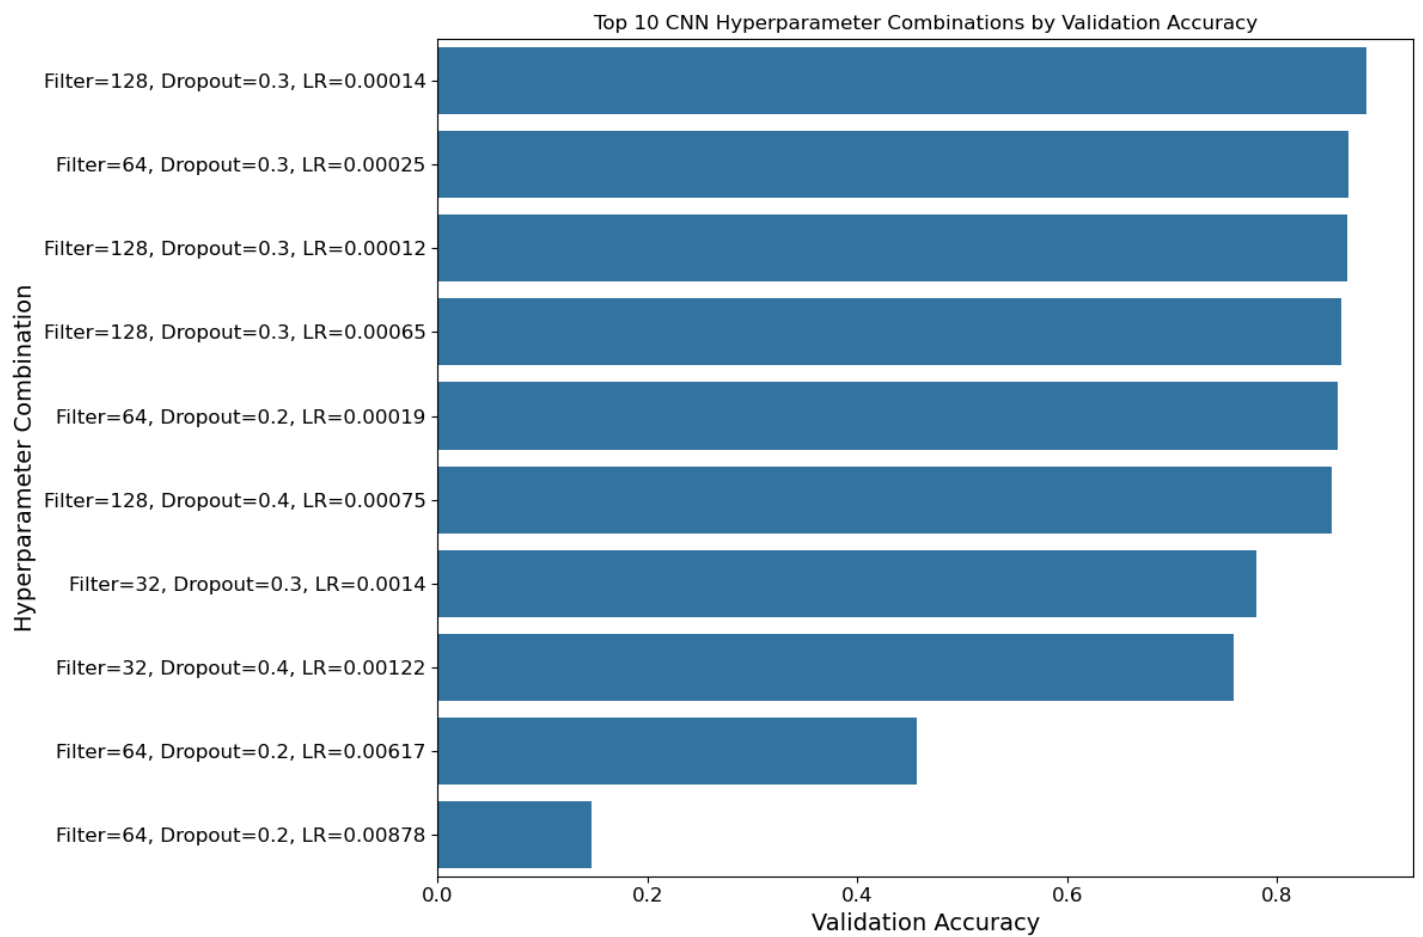
```

In [934... `# Sort by accuracy - DESC`  
`cnn_df_sorted = cnn_df.sort_values('score', ascending=False)`


```

# barplot (top 10 hyperparameter combinations)
plt.figure(figsize=(12, 8))
sns.barplot(x='score', y='param_combo', data=cnn_df_sorted.head(10))
plt.xlabel('Validation Accuracy')
plt.ylabel('Hyperparameter Combination')
plt.title('Top 10 CNN Hyperparameter Combinations by Validation Accuracy')
plt.tight_layout()
plt.show()

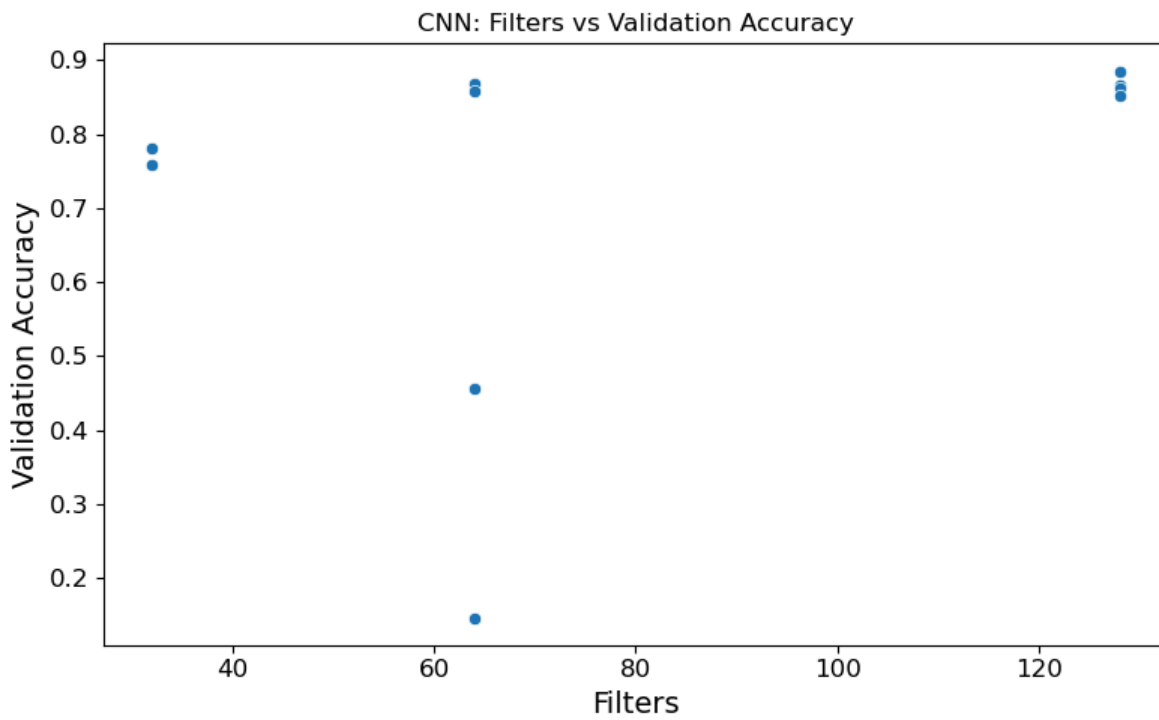
```




```
In [936... # Dropout Rate vs Validation Accuracy (bar)
plt.figure(figsize=(8, 5))
sns.barplot(x='dropout_rate', y='score', data=cnn_df)
plt.title('CNN: Dropout Rate vs Validation Accuracy')
plt.xlabel('Dropout Rate')
plt.ylabel('Validation Accuracy')
plt.tight_layout()
plt.show()
```



```
In [942... # Dense Units vs Validation Accuracy (scatter)
plt.figure(figsize=(8, 5))
sns.scatterplot(x='filters', y='score', data=cnn_df)
plt.title('CNN: Filters vs Validation Accuracy')
plt.xlabel('Filters')
plt.ylabel('Validation Accuracy')
plt.tight_layout()
plt.show()
```



```
In [944... # Learning Rate vs Validation Accuracy (log scale)
plt.figure(figsize=(8, 5))
sns.scatterplot(x='lr', y='score', data=cnn_df)
plt.xscale('log')
plt.title('CNN: Learning Rate vs Validation Accuracy')
plt.xlabel('Learning Rate (log scale)')
plt.ylabel('Validation Accuracy')
plt.tight_layout()
plt.show()
```



## Validation Model with Best Hyperparameters

To evaluate the training dynamics of the CNN model using the best-performing hyperparameter configuration, we manually rebuild the model using:

- filters = 128
- dropout\_rate = 0.3
- learning\_rate = 0.000136

The model includes three convolutional layers with increasing depth, each followed by pooling and dropout layers (except the last conv layer). A fully connected layer is added before the final classification layer.

We train the model for a fixed number of epochs = 30, and plot the accuracy and loss curves on both training and validation sets.

```
In [102... def valid_cnn():
    model = keras.models.Sequential()
    model.add(keras.layers.Input(shape = (28, 28, 3)))

    # Fixed parameters
    # Block 1
    model.add(keras.layers.Conv2D(filters = 128, kernel_size = (3, 3), activation = 'relu'))
    model.add(keras.layers.MaxPooling2D(pool_size = (2, 2)))
    model.add(keras.layers.Dropout(0.3))

    # Block 2
    model.add(keras.layers.Conv2D(filters = 128 * 2, kernel_size = (3, 3), activation = 'relu'))
    model.add(keras.layers.MaxPooling2D(pool_size = (2, 2)))
    model.add(keras.layers.Dropout(0.3))

    # Block 3
    model.add(keras.layers.Conv2D(filters = 128 * 2, kernel_size = (3, 3), activation = 'relu'))

    model.add(keras.layers.Flatten())

    model.add(keras.layers.Dense(128 * 2, activation='relu'))
    model.add(keras.layers.Dropout(0.3))

    model.add(keras.layers.Dense(9, activation = 'softmax'))

    model.compile(
        optimizer = keras.optimizers.Adam(learning_rate = 0.000136),
        loss = 'sparse_categorical_crossentropy',
        metrics = ['accuracy']
    )

    return model
```

```
In [102... # Fixed epoch =28
valid = valid_cnn()
history_cnn = valid.fit(
    X_train_nn, y_train_nn,
    epochs = 30,
    validation_data = (X_valid, y_valid)
)

plt.figure(figsize=(12, 4))


# Accuracy
plt.subplot(1, 2, 1)
plt.plot(history_cnn.history['accuracy'], label='Train Accuracy')
plt.plot(history_cnn.history['val_accuracy'], label='Val Accuracy')
plt.title('CNN - Training & Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

# Loss
plt.subplot(1, 2, 2)
plt.plot(history_cnn.history['loss'], label='Train Loss')
plt.plot(history_cnn.history['val_loss'], label='Val Loss')
plt.title('CNN - Training & Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```



Epoch 1/30  
**900/900** ————— 32s 34ms/step - accuracy: 0.2970 - loss: 1.8241 - val\_accuracy: 0.4897 - val\_loss: 1.3582  
Epoch 2/30  
**900/900** ————— 31s 34ms/step - accuracy: 0.5505 - loss: 1.2006 - val\_accuracy: 0.5981 - val\_loss: 1.0841  
Epoch 3/30  
**900/900** ————— 40s 44ms/step - accuracy: 0.6440 - loss: 0.9755 - val\_accuracy: 0.6913 - val\_loss: 0.8380  
Epoch 4/30  
**900/900** ————— 35s 39ms/step - accuracy: 0.6916 - loss: 0.8411 - val\_accuracy: 0.7422 - val\_loss: 0.7129  
Epoch 5/30  
**900/900** ————— 32s 35ms/step - accuracy: 0.7188 - loss: 0.7702 - val\_accuracy: 0.7541 - val\_loss: 0.6684  
Epoch 6/30  
**900/900** ————— 34s 38ms/step - accuracy: 0.7370 - loss: 0.7215 - val\_accuracy: 0.7572 - val\_loss: 0.6447  
Epoch 7/30  
**900/900** ————— 34s 37ms/step - accuracy: 0.7526 - loss: 0.6804 - val\_accuracy: 0.7734 - val\_loss: 0.6152  
Epoch 8/30  
**900/900** ————— 32s 36ms/step - accuracy: 0.7650 - loss: 0.6432 - val\_accuracy: 0.7956 - val\_loss: 0.5592  
Epoch 9/30  
**900/900** ————— 31s 34ms/step - accuracy: 0.7758 - loss: 0.6139 - val\_accuracy: 0.7966 - val\_loss: 0.5432  
Epoch 10/30  
**900/900** ————— 33s 36ms/step - accuracy: 0.7844 - loss: 0.5907 - val\_accuracy: 0.8122 - val\_loss: 0.5102  
Epoch 11/30  
**900/900** ————— 31s 34ms/step - accuracy: 0.7957 - loss: 0.5628 - val\_accuracy: 0.8159 - val\_loss: 0.4893  
Epoch 12/30  
**900/900** ————— 31s 34ms/step - accuracy: 0.8041 - loss: 0.5316 - val\_accuracy: 0.8256 - val\_loss: 0.4695  
Epoch 13/30  
**900/900** ————— 31s 35ms/step - accuracy: 0.8162 - loss: 0.5025 - val\_accuracy: 0.8344 - val\_loss: 0.4584  
Epoch 14/30  
**900/900** ————— 31s 34ms/step - accuracy: 0.8218 - loss: 0.4845 - val\_accuracy: 0.8369 - val\_loss: 0.4360  
Epoch 15/30  
**900/900** ————— 33s 36ms/step - accuracy: 0.8286 - loss: 0.4597 - val\_accuracy: 0.8406 - val\_loss: 0.4296  
Epoch 16/30  
**900/900** ————— 33s 37ms/step - accuracy: 0.8388 - loss: 0.4419 - val\_accuracy: 0.8456 - val\_loss: 0.4208  
Epoch 17/30  
**900/900** ————— 31s 34ms/step - accuracy: 0.8476 - loss: 0.4157 - val\_accuracy: 0.8416 - val\_loss: 0.4198  
Epoch 18/30  
**900/900** ————— 32s 35ms/step - accuracy: 0.8528 - loss: 0.4021 - val\_accuracy: 0.8356 - val\_loss: 0.4416  
Epoch 19/30  
**900/900** ————— 30s 34ms/step - accuracy: 0.8584 - loss: 0.3856 - val\_accuracy: 0.8378 - val\_loss: 0.4327  
Epoch 20/30  
**900/900** ————— 32s 36ms/step - accuracy: 0.8683 - loss: 0.3692 - val\_accuracy: 0.8447 - val\_loss: 0.4216  
Epoch 21/30  
**900/900** ————— 34s 38ms/step - accuracy: 0.8700 - loss: 0.3563 - val\_accuracy: 0.8497 - val\_loss: 0.4186  
Epoch 22/30  
**900/900** ————— 35s 39ms/step - accuracy: 0.8754 - loss: 0.3450 - val\_accuracy: 0.8516 - val\_loss: 0.4115  
Epoch 23/30  
**900/900** ————— 32s 36ms/step - accuracy: 0.8819 - loss: 0.3261 - val\_accuracy: 0.8444 - val\_loss: 0.4168  
Epoch 24/30  
**900/900** ————— 32s 36ms/step - accuracy: 0.8837 - loss: 0.3205 - val\_accuracy: 0.8406 - val\_loss: 0.4308  
Epoch 25/30  
**900/900** ————— 31s 35ms/step - accuracy: 0.8855 - loss: 0.3102 - val\_accuracy: 0.8375 - val\_loss: 0.4437  
Epoch 26/30  
**900/900** ————— 32s 35ms/step - accuracy: 0.8901 - loss: 0.2973 - val\_accuracy: 0.8591 - val\_loss: 0.3830  
Epoch 27/30  
**900/900** ————— 32s 35ms/step - accuracy: 0.8974 - loss: 0.2837 - val\_accuracy: 0.8572 - val\_loss: 0.3927  
Epoch 28/30  
**900/900** ————— 32s 36ms/step - accuracy: 0.9010 - loss: 0.2729 - val\_accuracy: 0.8525 - val\_loss: 0.4025  
Epoch 29/30  
**900/900** ————— 32s 36ms/step - accuracy: 0.9029 - loss: 0.2729 - val\_accuracy: 0.8641 - val\_loss: 0.3930  
Epoch 30/30  
**900/900** ————— 31s 35ms/step - accuracy: 0.9088 - loss: 0.2494 - val\_accuracy: 0.8334 - val\_loss: 0.4660



We did not choose the best epoch = 30 returned by the parameter search. Instead, we selected epoch = 18 for the final model training, for the following reasons:

- **Validation loss (val\_loss) reached its lowest point around epoch 15**
  - val\_loss reached a minimum at around epoch 18, then started to slightly increase, which may indicate early signs of overfitting.

- **Validation accuracy (val\_accuracy) also became stable or slightly decreased after epoch 18**
  - As shown in the left plot, val\_accuracy peaked around epoch 15, then fluctuated without further improvement;
  - Meanwhile, train\_accuracy kept increasing, showing that the model was still fitting the training set and possibly starting to overfit.

Therefore, based on the plots above, we consider **epoch = 18** to be a good balance between generalization and training performance, and we use it as the final training setting.

## 4. Final models

In this section, we train each model with its **best hyperparameter combination independently of the hyperparameter tuning cells**, don't rely on the hyperparameter tuning cells having been run.

### Algorithm of choice from first six weeks of course - Random Forest

```
In [226... # Rebuild a new RF based on the optimal parameters
final_rf = RandomForestClassifier(
    n_estimators = 200,
    max_depth = 30,
    min_samples_split = 5,
    max_features = 'sqrt',
    random_state=42
)

# Retrain on the training set
final_rf.fit(X_train_rf, y_train)
```

```
Out [226... RandomForestClassifier

RandomForestClassifier(max_depth=30, min_samples_split=5, n_estimators=200,
                      random_state=42)
```

```
In [149... # Evaluate on the test set
y_pred_test_rf = final_rf.predict(X_test_rf)
report_rf = classification_report(y_test, y_pred_test_rf, output_dict=True)

acc_rf = report_rf['accuracy']
macro_rf = report_rf['macro avg']['f1-score']


print("Final RF - Test Set Classification Report:")
print(classification_report(y_test, y_pred_test_rf))
```

```
Final RF - Test Set Classification Report:
              precision    recall  f1-score   support

    0           0.81       0.97       0.89         873
    1           0.99       0.90       0.94         858
    2           0.66       0.31       0.43         877
    3           0.78       0.78       0.78         914
    4           0.54       0.66       0.60         737
    5           0.64       0.73       0.69        1072
    6           0.55       0.12       0.19         682
    7           0.51       0.49       0.50         813
    8           0.51       0.78       0.62        1174

 accuracy          0.66         0.66         0.66        8000
 macro avg         0.67         0.64         0.63        8000
 weighted avg         0.67         0.66         0.64        8000
```

```
In [150... # Print the confusion matrix
cm = confusion_matrix(y_test, y_pred_test_rf)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Random Forest")
plt.show()
```



## Fully connected neural network

```
In [171... # Rebuild a new MLP based on the optimal parameters
def build_final_mlp():
    model = keras.models.Sequential()
    model.add(keras.layers.Input(shape = (28, 28, 3)))
    model.add(keras.layers.Flatten())

    # Fixed parameter
    model.add(keras.layers.Dense(448, activation = 'relu'))
    model.add(keras.layers.Dropout(0.2))

    model.add(keras.layers.Dense(320, activation = 'relu'))
    model.add(keras.layers.Dropout(0.2))

    model.add(keras.layers.Dense(9, activation = 'softmax'))

    model.compile(
        optimizer = keras.optimizers.Adam(learning_rate = 0.000104),
        loss = 'sparse_categorical_crossentropy',
        metrics = ['accuracy']
    )

    return model
```

Previously, we used `train_test_split()` to allocate **90% of the training data for hyperparameter tuning and 10% of the training data for validation**.

- Now that the optimal hyperparameters have been selected, we will train the final model using the **full training set**, without using `X_valid`, since it has already been involved in the tuning process.

```
In [174... final_mlp = build_final_mlp()

# Retrain on the overall training set
history_mlp = final_mlp.fit(X_train_scale,
                             y_train,
                             epochs = 28
)
```

```

Epoch 1/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.2210 - loss: 2.0389
Epoch 2/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.3268 - loss: 1.7498
Epoch 3/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.3718 - loss: 1.6659
Epoch 4/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.3947 - loss: 1.6083
Epoch 5/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4162 - loss: 1.5564
Epoch 6/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4364 - loss: 1.5225
Epoch 7/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4452 - loss: 1.4877
Epoch 8/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4645 - loss: 1.4568
Epoch 9/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4626 - loss: 1.4424
Epoch 10/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4732 - loss: 1.4234
Epoch 11/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4728 - loss: 1.4115
Epoch 12/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4827 - loss: 1.4009
Epoch 13/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4860 - loss: 1.3812
Epoch 14/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4890 - loss: 1.3751
Epoch 15/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4972 - loss: 1.3603
Epoch 16/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.4970 - loss: 1.3534
Epoch 17/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5011 - loss: 1.3430
Epoch 18/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5002 - loss: 1.3346
Epoch 19/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5042 - loss: 1.3276
Epoch 20/28
1000/1000 ————— 5s 5ms/step - accuracy: 0.5052 - loss: 1.3166
Epoch 21/28
1000/1000 ————— 5s 5ms/step - accuracy: 0.5059 - loss: 1.3153
Epoch 22/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5138 - loss: 1.2986
Epoch 23/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5163 - loss: 1.2926
Epoch 24/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5223 - loss: 1.2889
Epoch 25/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5224 - loss: 1.2785
Epoch 26/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5201 - loss: 1.2750
Epoch 27/28
1000/1000 ————— 4s 4ms/step - accuracy: 0.5247 - loss: 1.2614
Epoch 28/28
1000/1000 ————— 5s 5ms/step - accuracy: 0.5268 - loss: 1.2578

```

```
In [175]: final_mlp.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 2352)	0
dense_8 (Dense)	(None, 448)	1,054,144
dropout_2 (Dropout)	(None, 448)	0
dense_9 (Dense)	(None, 320)	143,680
dropout_3 (Dropout)	(None, 320)	0
dense_10 (Dense)	(None, 9)	2,889

Total params: 3,602,141 (13.74 MB)

Trainable params: 1,200,713 (4.58 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 2,401,428 (9.16 MB)

```
final_mlp.evaluate(X_test_scale, y_test)
```

- It calculates the **overall accuracy** of the model.

```
In [177... # Evaluate on the test set
loss_mlp, acc_mlp = final_mlp.evaluate(X_test_scale, y_test)
print(f"Accuracy on test data: {acc_mlp:.4f}")
```

250/250 ————— 1s 1ms/step - accuracy: 0.5415 - loss: 1.1903  
Accuracy on test data: 0.5490

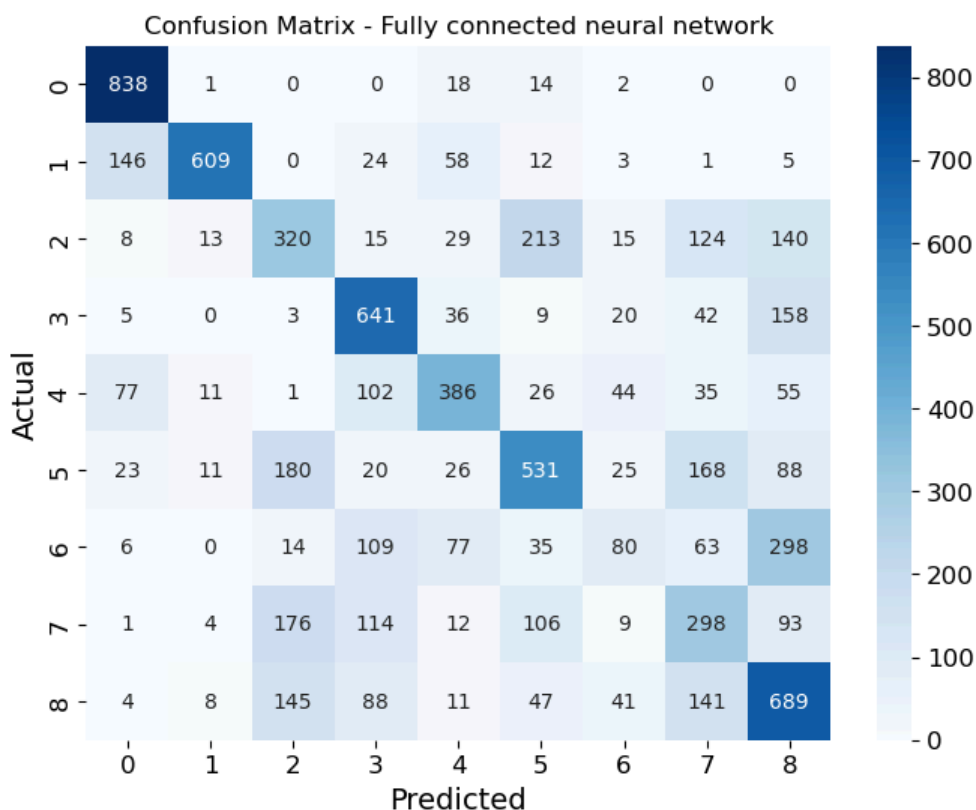
```
In [178... y_pred_mlp = np.argmax(final_mlp.predict(X_test_scale), axis=1)
report_mlp = classification_report(y_test, y_pred_mlp, output_dict=True)
macro_mlp = report_mlp['macro avg']['f1-score']

print(classification_report(y_test, y_pred_mlp))
```

250/250 ————— 0s 889us/step

	precision	recall	f1-score	support
0	0.76	0.96	0.85	873
1	0.93	0.71	0.80	858
2	0.38	0.36	0.37	877
3	0.58	0.70	0.63	914
4	0.59	0.52	0.56	737
5	0.53	0.50	0.51	1072
6	0.33	0.12	0.17	682
7	0.34	0.37	0.35	813
8	0.45	0.59	0.51	1174
accuracy			0.55	8000
macro avg	0.54	0.54	0.53	8000
weighted avg	0.55	0.55	0.54	8000

```
In [179... # Print the confusion matrix
cm = confusion_matrix(y_test, y_pred_mlp)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Fully connected neural network")
plt.show()
```



## Convolutional neural network

```
In [186... # Rebuild a new CNN based on the optimal parameters
def build_final_cnn():
    model = keras.models.Sequential()
    model.add(keras.layers.Input(shape = (28, 28, 3)))

    # Block 1
    model.add(keras.layers.Conv2D(filters = 128, kernel_size = (3, 3), activation = 'relu'))
    model.add(keras.layers.MaxPooling2D(pool_size = (2, 2)))
    model.add(keras.layers.Dropout(0.3))
```

```

# Block 2
model.add(keras.layers.Conv2D(filters = 128 * 2, kernel_size = (3, 3), activation = 'relu'))
model.add(keras.layers.MaxPooling2D(pool_size = (2, 2)))
model.add(keras.layers.Dropout(0.3))

model.add(keras.layers.Conv2D(filters = 128 * 2, kernel_size = (3, 3), activation = 'relu'))

model.add(keras.layers.Flatten())

model.add(keras.layers.Dense(128 * 2, activation='relu'))
model.add(keras.layers.Dropout(0.3))

model.add(keras.layers.Dense(9, activation = 'softmax'))

model.compile(
    optimizer = keras.optimizers.Adam(learning_rate = 0.000136),
    loss = 'sparse_categorical_crossentropy',
    metrics = ['accuracy']
)

return model

```

```

In [188... # epoch = 18 (not 30 -> reduce computing time and the performance is not much different)
final_cnn = build_final_cnn()

# Retrain on the overall training set
history_cnn = final_cnn.fit(
    X_train_scale, y_train,
    epochs = 18
)

```

```

Epoch 1/18
1000/1000 ————— 30s 29ms/step - accuracy: 0.2993 - loss: 1.8284
Epoch 2/18
1000/1000 ————— 32s 32ms/step - accuracy: 0.5642 - loss: 1.1567
Epoch 3/18
1000/1000 ————— 32s 32ms/step - accuracy: 0.6474 - loss: 0.9617
Epoch 4/18
1000/1000 ————— 33s 33ms/step - accuracy: 0.6934 - loss: 0.8410
Epoch 5/18
1000/1000 ————— 35s 35ms/step - accuracy: 0.7240 - loss: 0.7568
Epoch 6/18
1000/1000 ————— 34s 34ms/step - accuracy: 0.7448 - loss: 0.6991
Epoch 7/18
1000/1000 ————— 33s 33ms/step - accuracy: 0.7631 - loss: 0.6540
Epoch 8/18
1000/1000 ————— 34s 34ms/step - accuracy: 0.7736 - loss: 0.6283
Epoch 9/18
1000/1000 ————— 35s 35ms/step - accuracy: 0.7846 - loss: 0.5981
Epoch 10/18
1000/1000 ————— 34s 34ms/step - accuracy: 0.7966 - loss: 0.5674
Epoch 11/18
1000/1000 ————— 33s 33ms/step - accuracy: 0.8024 - loss: 0.5500
Epoch 12/18
1000/1000 ————— 31s 31ms/step - accuracy: 0.8151 - loss: 0.5107
Epoch 13/18
1000/1000 ————— 29s 29ms/step - accuracy: 0.8208 - loss: 0.4980
Epoch 14/18
1000/1000 ————— 29s 29ms/step - accuracy: 0.8312 - loss: 0.4668
Epoch 15/18
1000/1000 ————— 29s 29ms/step - accuracy: 0.8325 - loss: 0.4536
Epoch 16/18
1000/1000 ————— 29s 29ms/step - accuracy: 0.8445 - loss: 0.4279
Epoch 17/18
1000/1000 ————— 30s 30ms/step - accuracy: 0.8484 - loss: 0.4121
Epoch 18/18
1000/1000 ————— 29s 29ms/step - accuracy: 0.8545 - loss: 0.4039

```

```

In [189... final_cnn.summary()

```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 128)	3,584
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 128)	0
dropout_4 (Dropout)	(None, 13, 13, 128)	0
conv2d_3 (Conv2D)	(None, 11, 11, 256)	295,168
max_pooling2d_3 (MaxPooling2D)	(None, 5, 5, 256)	0
dropout_5 (Dropout)	(None, 5, 5, 256)	0
conv2d_4 (Conv2D)	(None, 3, 3, 256)	590,080
flatten_4 (Flatten)	(None, 2304)	0
dense_11 (Dense)	(None, 256)	590,080
dropout_6 (Dropout)	(None, 256)	0
dense_12 (Dense)	(None, 9)	2,313

**Total params:** 4,443,677 (16.95 MB)

**Trainable params:** 1,481,225 (5.65 MB)

**Non-trainable params:** 0 (0.00 B)

**Optimizer params:** 2,962,452 (11.30 MB)

```
In [190... # Evaluate on the test set
loss_cnn, acc_cnn = final_cnn.evaluate(X_test_scale, y_test)
print(f"Accuracy on test data: {acc_cnn:.4f}")

250/250 ————— 2s 10ms/step - accuracy: 0.8557 - loss: 0.3834
Accuracy on test data: 0.8537
```

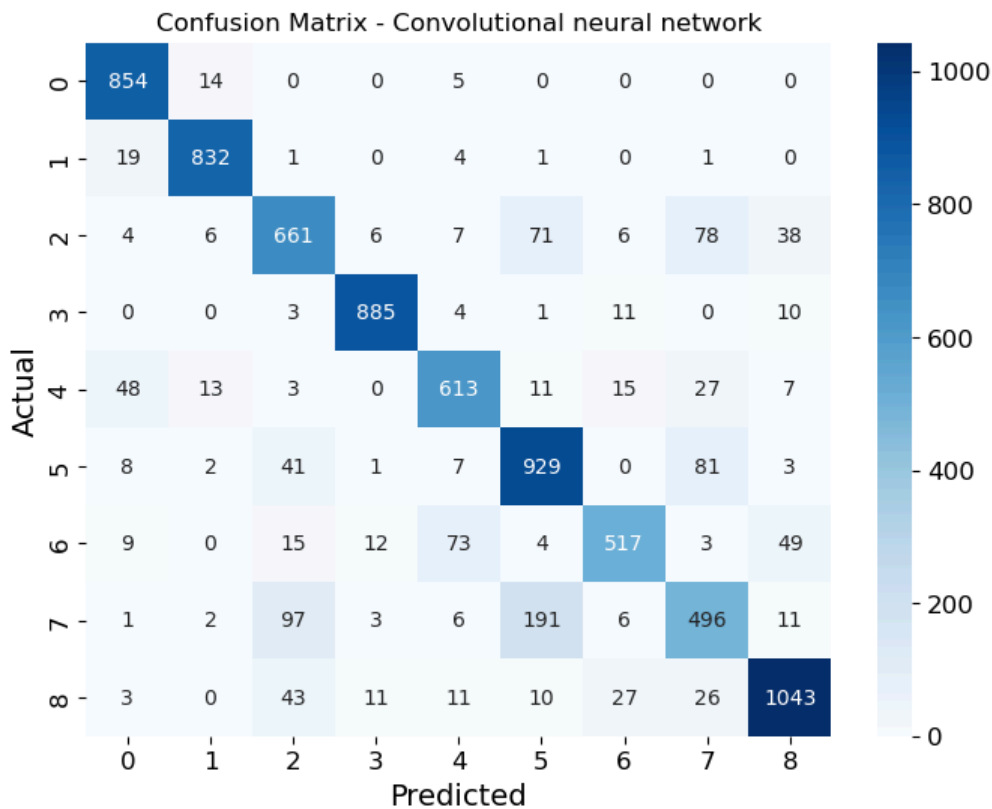
```
In [197... y_pred_cnn = np.argmax(final_cnn.predict(X_test_scale), axis=1)
report_cnn = classification_report(y_test, y_pred_cnn, output_dict=True)
macro_cnn = report_cnn['macro avg']['f1-score']

print(classification_report(y_test, y_pred_cnn))

250/250 ————— 2s 9ms/step
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	873
1	0.96	0.97	0.96	858
2	0.77	0.75	0.76	877
3	0.96	0.97	0.97	914
4	0.84	0.83	0.84	737
5	0.76	0.87	0.81	1072
6	0.89	0.76	0.82	682
7	0.70	0.61	0.65	813
8	0.90	0.89	0.89	1174
accuracy			0.85	8000
macro avg	0.85	0.85	0.85	8000
weighted avg	0.85	0.85	0.85	8000

```
In [199... # confusion matrix
cm = confusion_matrix(y_test, y_pred_cnn)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Convolutional neural network")
plt.show()
```



## Model comparison

- **Accuracy** is sensitive to class imbalance and may mask poor performance on minority classes.
- **Macro F1** averages the F1 scores across all classes without considering class weights. It is less affected by class distribution and better reflects overall fairness.

```
In [234... models = ['Random Forest', 'MLP', 'CNN']
accuracies = [acc_rf, acc_mlp, acc_cnn]
f1_scores = [macro_rf, macro_mlp, macro_cnn]

x = np.arange(len(models))
width = 0.35

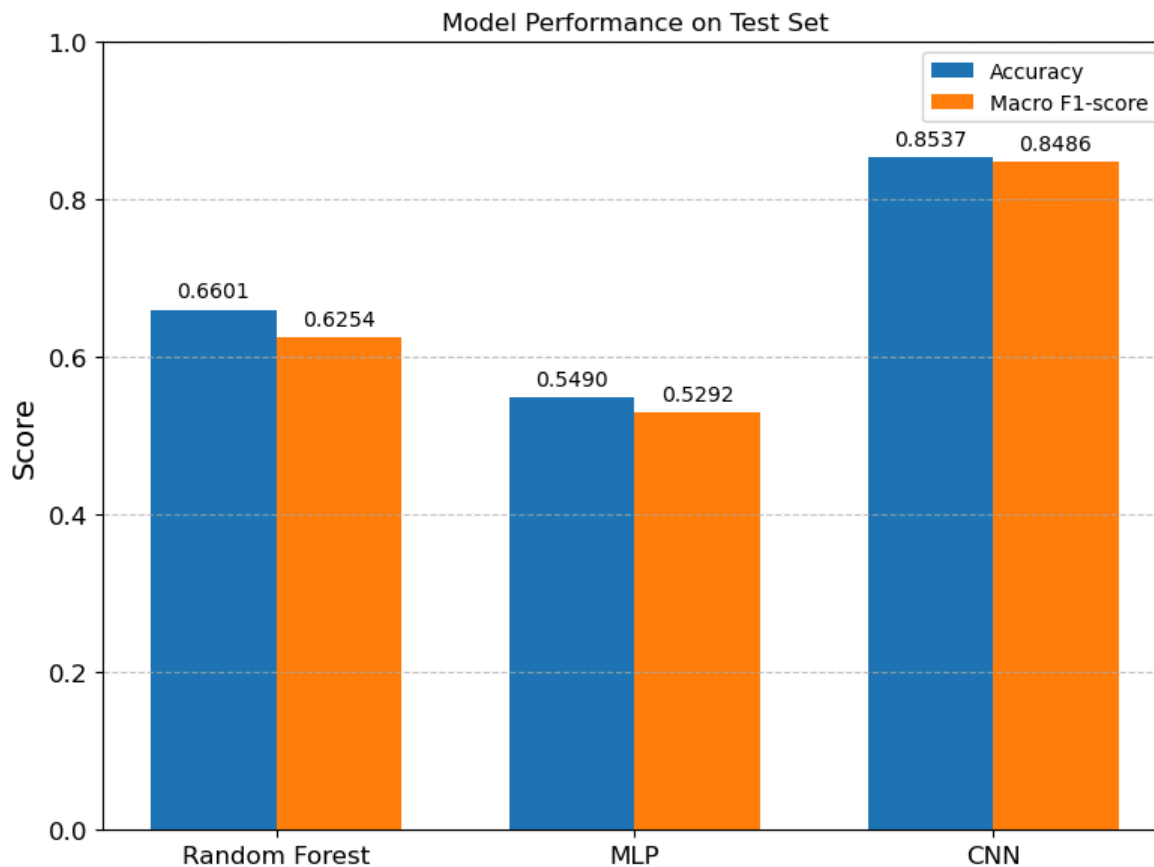
# Model Performance on Test Set
plt.figure(figsize=(8, 6))
bars1 = plt.bar(x - width/2, accuracies, width, label='Accuracy')
bars2 = plt.bar(x + width/2, f1_scores, width, label='Macro F1-score')

plt.xticks(x, models)
plt.ylim(0, 1)
plt.ylabel('Score')
plt.title('Model Performance on Test Set')
plt.legend()
plt.grid(True, axis='y', linestyle='--', alpha=0.7)

# Display the value on each bar
for bar in bars1 + bars2:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2,
             height + 0.01,
             f'{height:.4f}',
             ha='center', va='bottom')

plt.tight_layout()
plt.show()
```





**Model performance on the validation set:** Since we do not repeat hyperparameter tuning, the final testing phase uses the full training and test sets. Therefore, the validation scores shown here refer to the best validation performance recorded during the tuning process under the optimal hyperparameter combination.


```
In [221]: # Validation vs Test Accuracy Comparison
val_scores = [0.649875, 0.551250, 0.885312]

plt.figure(figsize=(8, 5))
bars3 = plt.bar(x - width/2, val_scores, width, label='Validation Accuracy')
bars4 = plt.bar(x + width/2, accuracies, width, label='Test Accuracy')

plt.xticks(x, models)
plt.ylim(0, 1)
plt.ylabel('Accuracy')
plt.title('Validation vs Test Accuracy Comparison')
plt.legend()
plt.grid(True, axis='y', linestyle='--', alpha=0.7)

# Display the value on each bar
for bar in bars3 + bars4:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2,
             height + 0.01,
             f'{height:.4f}',
             ha='center', va='bottom')

plt.tight_layout()
plt.show()
```



**Baseline model vs Final model:** The figure compares the accuracy performance of the three models on the test set after baseline and hyperparameter tuning.

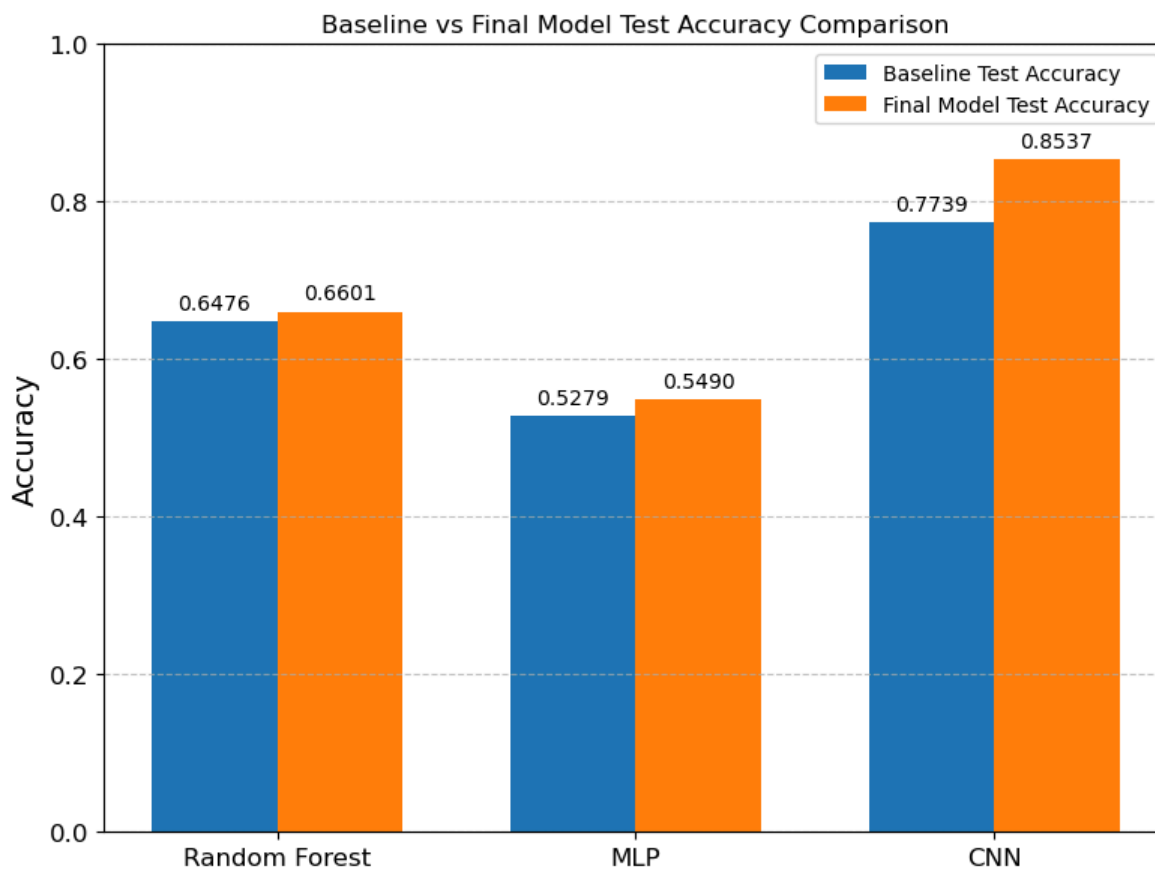
```
In [224... # Baseline vs Final Model Test Accuracy Comparison
base_scores = [acc_base_rf, acc_base_mlp, acc_base_cnn]

plt.figure(figsize=(8, 6))
bars5 = plt.bar(x - width/2, base_scores, width, label='Baseline Test Accuracy')
bars6 = plt.bar(x + width/2, accuracies, width, label='Final Model Test Accuracy')

plt.xticks(x, models)
plt.ylim(0, 1)
plt.ylabel('Accuracy')
plt.title('Baseline vs Final Model Test Accuracy Comparison')
plt.legend()
plt.grid(True, axis='y', linestyle='--', alpha=0.7)

# Display the value on each bar
for bar in bars5 + bars6:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2,
             height + 0.01,
             f'{height:.4f}',
             ha='center', va='bottom')

plt.tight_layout()
plt.show()
```



```
In [266... # Hyperparameter tuning time (Unit: minutes)
tuning_times = [100, 22, 73]

plt.figure(figsize=(10, 6.5))


bars = plt.bar(models, tuning_times, color=['green', 'skyblue', 'red'])

for bar in bars:
    height = bar.get_height()
    plt.text(
        bar.get_x() + bar.get_width()/2,
        height + 0.01,
        f'{height} min',
        ha='center', va='bottom')

plt.title('Hyperparameter Tuning Time Comparison (Unit: minutes)')
plt.ylabel('Tuning Time (minutes)')
plt.ylim(0, max(tuning_times) + 3)
plt.grid(axis='y', linestyle='--', alpha=0.7)

plt.tight_layout()
plt.show()
```

Hyperparameter Tuning Time Comparison (Unit: minutes)



In [ ]: